


DUAL-PATH W4A16 GEMM VS CUBLAS FP16: EMPIRICAL REGIME BOUNDARIES

Speedup across Batch Sizes $M \in \{1, 4, 8, 16, 32, 64\}$ on Physical Tesla T4 Silicon (Table 4)

Panel A: Speedup Factor vs Batch Size M

CP-Hybrid GEMV ($M \leq 4$)



Panel B: Execution Latency at Decode ($M=1$)

